

ALMA MATER STUDIORUM – UNIVERSITÀ DI BOLOGNA

School of Engineering and Architecture

Master's Degree in Artificial Intelligence

Machine Learning for Computer Vision (ML4CV)  
3 CFU Project

# **Improving Real-Time Cone Detection for Autonomous Formula SAE Racing with Modern YOLO Architectures**

**Student:**

Nicolas Cridlig

**Supervisor:**

Prof. Samuele Salti

Academic Year 2024–2025

January 2026

# **Abstract**

# Contents

<b>Abstract</b>	<b>1</b>
<b>1 Introduction</b>	<b>6</b>
1.1 Context: Formula Student Driverless . . . . .	6
1.2 Problem Statement . . . . .	6
1.3 Objectives . . . . .	7
1.4 Contributions . . . . .	7
1.5 Report Structure . . . . .	8
<b>2 Background</b>	<b>9</b>
2.1 Object Detection with YOLO . . . . .	9
2.1.1 YOLOv11 . . . . .	9
2.1.2 YOLOv12 . . . . .	9
2.1.3 YOLOv26 . . . . .	10
2.2 Transfer Learning and Catastrophic Forgetting . . . . .	10
2.3 Model Optimization for Edge Deployment . . . . .	10
2.4 Formula Student Driverless Perception . . . . .	11
2.5 Related Work . . . . .	11
<b>3 Methods</b>	<b>13</b>
3.1 Experimental Setup . . . . .	13
3.2 Datasets . . . . .	13
3.2.1 FSOCO-12 . . . . .	14
3.2.2 cone-detector . . . . .	14
3.2.3 fsoco-ubm: Real-World Test Set . . . . .	15
3.3 Baseline Reproduction . . . . .	15
3.4 Hyperparameter Sweep . . . . .	16
3.5 Architecture Evaluation . . . . .	16
3.6 Two-Stage Training Strategy . . . . .	17
3.7 Deployment Pipeline . . . . .	17

<b>4 Results</b>	<b>19</b>
4.1 FSOCO-12 Benchmark Results . . . . .	19
4.2 Per-Class Analysis . . . . .	20
4.3 Two-Stage vs. Single-Stage Training . . . . .	20
4.4 Hyperparameter Sweep . . . . .	21
4.5 Real-World Validation (fsoco-ubm) . . . . .	22
4.6 Deployment Performance . . . . .	22
<b>5 Conclusion</b>	<b>24</b>
5.1 Summary of Contributions . . . . .	24
5.2 Lessons Learned . . . . .	24
5.3 Limitations . . . . .	25
5.4 Future Work . . . . .	25
<b>A Technical Details</b>	<b>27</b>
A.1 Hyperparameter Sweep Configuration . . . . .	27
A.2 Two-Stage Training Hyperparameters . . . . .	27
A.3 The <code>optimizer='auto'</code> Bug . . . . .	27
A.4 fsoco-ubm Dataset Creation Pipeline . . . . .	27

# List of Figures

1.1	Example false positive: a person misclassified as a yellow cone by the baseline YOLOv11n detector. . . . .	7
3.1	Class distribution in the FSOCO-12 dataset. . . . .	14
3.2	3D visualization of detected cones from the stereo pipeline. . . . .	18

# List of Tables

3.1	Hardware used for training and deployment. . . . .	13
3.2	Datasets used in this project. . . . .	14
3.3	Architecture specifications for the three evaluated models. . . . .	17
4.1	Overall performance on the FSOCO-12 test set (689 images, 12,054 instances). . . . .	19
4.2	Per-class performance of YOLO26n (single-stage) on FSOCO-12 test set. . . . .	20
4.3	Per-class performance of YOLO12n on FSOCO-12 test set. . . . .	20
4.4	Two-stage vs. single-stage YOLO26n training comparison. . . . .	21
4.5	Hyperparameter sweep summary (W&B Bayesian optimization). . . . .	21
4.6	Real-world performance on fsoco-ubm test set (96 images, 1,426 instances). . . . .	22
4.7	Inference latency breakdown on RTX 4060 (TensorRT FP16). . . . .	23
4.8	Deployment comparison: YOLO26n vs. UBM baseline. . . . .	23
A.1	Two-stage training configuration. . . . .	28

# Chapter 1

## Introduction

### 1.1 Context: Formula Student Driverless

This report addresses an element of the work undertaken by Unibo Motorsport (UBM) to succeed in the Formula SAE (FSAE) competition. Since 2017, the competition requires cars to autonomously navigate an unknown track delineated by colored traffic cones at maximum speed. To do so, the software and hardware stack must be efficient, reliable, and well integrated. The perception of the colored cones is vital due to the garbage-in, garbage-out moniker: even the best written code cannot produce good results on bad data. Therefore, in this report we increase successful cone detections by utilizing the newest generation of YOLO models alongside a new dataset. This research is vital, last summer, UBM competed at Formula Student Germany, and this year we will compete again from the 11-16th of August 2026.

### 1.2 Problem Statement

For perception, the software stack relies upon a custom trained YOLO object detector to identify track cones from a stereocamera so that they may be reconstructed with a multi-stage matching algorithm. The baseline model, YOLOv11, has room for improvement. It is trained on base hyperparameters and any tuning or analysis was lost. The  $mAP_{50}=0.666$  and while this figure is our quantitative benchmark, qualitative issues registered by the team include exposure sensitivity and false positives. The model is constrained by both accuracy and real time speed. The stereocamera operates at 60fps which allows for a maximum control loop time of 16.7 ms. Any time perception takes out of this hard limit, is less time for the decision and control algorithms to operate.



Figure 1.1: Example false positive: a person misclassified as a yellow cone by the baseline YOLOv11n detector.

## 1.3 Objectives

The project objectives are:

1. Reproduce and establish a reproducible baseline
2. Evaluate modern YOLO architectures (YOLO12n, YOLO26n)
3. Explore two-stage training (pre-training on larger dataset + fine-tuning)
4. Assess hyperparameter optimization via Bayesian sweep
5. Deploy optimized model on RTX 4060 via TensorRT
6. Create real-world validation dataset (fsoco-ubm)

## 1.4 Contributions

We systematically evaluated three YOLO architectures on the cone detection task and found that YOLO26n, the most recent and lightweight of the three, achieves the best results. On the FSOCO-12 benchmark it reaches  $mAP_{50}=0.763$ , a 14.6% improvement over the model currently running on the car. We also explored a two-stage training strategy, pre-training on a larger 22,725 image dataset before fine-tuning on FSOCO-12, which trades a small amount of benchmark precision for better recall and real-world generalization. A Bayesian hyperparameter sweep over 13 parameters showed that the Ultralytics defaults are already near-optimal, providing no measurable gain. To validate our results beyond internet benchmarks, we created fsoco-ubm, a 96-image test set extracted from the car's own camera recordings at the Rioveggio test track. Finally, we deployed the selected model via TensorRT FP16 on the onboard RTX 4060, achieving 2.63 ms inference latency with a 6.3× margin over the 60 fps real-time requirement.

## 1.5 Report Structure

Chapter 2 covers the relevant background on YOLO architectures, transfer learning, and deployment optimization. Chapter 3 describes the datasets, training procedures, and deployment pipeline. Chapter 4 presents the benchmark results, per-class analysis, real-world validation, and deployment performance. Chapter 5 summarizes our findings and outlines future work. Chapter A documents technical details including the sweep configuration, two-stage hyperparameters, and a bug encountered with the Ultralytics optimizer.

# Chapter 2

## Background

### 2.1 Object Detection with YOLO

YOLO (You Only Look Once) is a family of real-time object detection models first introduced by Redmon et al. (2016) and now maintained by Ultralytics (Jocher et al., 2023). Unlike two-stage detectors that first propose regions and then classify them, YOLO performs detection in a single forward pass through the network, which is what makes it fast enough for our use case. Over the past decade the architecture has evolved considerably: modern variants use anchor-free detection heads, feature pyramid networks for multi-scale prediction, and CSP (Cross Stage Partial) backbones for efficient feature extraction. We evaluated three successive generations of nano-sized YOLO models, each offering a different trade-off between accuracy and computational cost.

#### 2.1.1 YOLOv11

YOLOv11 was released in 2024 by Ultralytics (Ultralytics, 2024) and is the architecture currently deployed on the UBM race car. It introduces C3k2 backbone blocks, a Spatial Pyramid Pooling Fast (SPPF) layer, and C2PSA attention modules. The nano variant has 2.59M parameters and 6.4 GFLOPs. This is our baseline: every other model in this report is measured against it.

#### 2.1.2 YOLO12

YOLO12, published in January 2025 by Tian et al. (2025), replaces the CSP backbone with R-ELAN (Residual Efficient Layer Aggregation Network) and introduces an area-attention mechanism that captures global context without the quadratic cost of standard self-attention. The nano variant sits at 2.56M parameters and 6.3 GFLOPs, marginally smaller than YOLOv11n. We tested it as a potential drop-in upgrade since the parameter count and inference profile are nearly identical.

### 2.1.3 YOLO26

YOLO26 is the most recent release from Ultralytics at the time of writing (Ultralytics, 2025). Its nano variant is the smallest of the three we evaluated: 2.51M parameters and 5.8 GFLOPs. The reduction in both parameters and floating point operations translates directly to faster inference, which is a key priority for deploying the model. It also eliminates the non-maximum suppression (NMS) post-processing step, further reducing latency.

## 2.2 Transfer Learning and Catastrophic Forgetting

Transfer learning is the practice of pre-training a model on a large dataset and then fine-tuning it on a smaller, task-specific one. In our case, we pre-train on the 22,725 image cone-detector dataset (FSB Driverless, 2024) and then fine-tune on the 9,778 image FSOCO-12 dataset (FMDV, 2024). The idea is that the model learns general cone features during pre-training and refines them during fine-tuning.

The main risk with this approach is catastrophic forgetting (French, 1999): the model overwrites the useful features it learned during pre-training when exposed to the new dataset. In our first fine-tuning attempt, validation mAP<sub>50</sub> dropped from 0.7536 to 0.6278 in just two epochs—a 16.7% collapse—before slowly recovering over 38 epochs. The root cause was Ultralytics’ `optimizer='auto'` default, which silently ignored our specified learning rate of 0.001 and instead used a hardcoded 0.01 (10× too high for fine-tuning). The framework’s internal heuristic selects SGD with fixed hyperparameters whenever the total iteration count exceeds 10,000, regardless of user settings. We resolved this by explicitly setting `optimizer='AdamW'` (Loshchilov and Hutter, 2019), which decouples weight decay from the gradient update and respects the user-specified learning rate. We also adopted a two-phase strategy: first training only the detection head with frozen backbone layers, then unfreezing all layers with a 100× lower learning rate and 20% warmup. This eliminated the forgetting entirely. The full debugging process is documented in section A.3.

## 2.3 Model Optimization for Edge Deployment

The models are trained in PyTorch but deployed through a multi-stage optimization pipeline. First, the trained weights are exported to ONNX (Open Neural Network Exchange) (ONNX Community, 2019), a hardware-agnostic intermediate representation. Then, NVIDIA TensorRT (NVIDIA Corporation, 2024) compiles the ONNX graph into an optimized engine for the target GPU. TensorRT applies several automatic optimizations: kernel fusion combines sequences of operations (e.g. Conv+BatchNorm+ReLU) into a

single GPU kernel to reduce memory bandwidth and kernel launch overhead, and precision reduction converts FP32 weights to FP16 or INT8 to exploit the GPU’s dedicated Tensor Cores. INT8 quantization requires a calibration dataset of 500–1,000 representative images to determine the optimal scaling factors, and typically incurs a 2–3% mAP<sub>50</sub> drop compared to FP16.

One important constraint is that TensorRT engines are tied to the specific GPU architecture they are built on. An engine compiled for the RTX 4080 Super (training machine) will not run on the RTX 4060 (race car). We therefore build the final engine directly on the car’s ASU (Autonomous System Unit). In our case, FP16 precision was sufficient: the resulting 2.63 ms inference time leaves a 6.3× margin over the 60 fps requirement, making INT8 quantization unnecessary and its accuracy penalty avoidable.

## 2.4 Formula Student Driverless Perception

The perception pipeline, as designed by [Fusa \(2025\)](#), is built around a ZED 2i stereo camera ([Stereolabs, 2021](#)) operating at 1280×720 resolution and 30 fps but with the capability of operating at 60 fps, awaiting testing. The YOLO detector runs on each frame to produce 2D bounding boxes for five cone classes defined by the FSAE Driverless rules ([SAE International, 2025](#)): small blue cones mark the left track border, small yellow cones mark the right border, small orange cones delineate entry and exit lanes, and large orange cones are placed before and after start, finish, and timekeeping lines. A fifth class, “unknown,” is used for ambiguous or occluded cones that cannot be confidently classified. These detections are then matched across the stereo pair and triangulated to produce 3D positions relative to the car. The entire pipeline must complete within 16.7 ms per frame to maintain the 60 fps loop rate; any excess time taken by detection directly reduces the budget available for planning and control downstream.

## 2.5 Related Work

The FSOCO (Formula Student Objects in Context) dataset ([FSOCO Contributors, 2020](#)) is a community-maintained benchmark for cone detection, contributed by multiple Formula Student teams. It provides a shared evaluation standard, though the images come from a variety of cameras and conditions that may not match any single team’s deployment setup. Most teams in the Driverless competition use some variant of YOLO, trained on FSOCO or private datasets, and evaluate using the standard COCO metrics ([Lin et al., 2014](#)): mAP<sub>50</sub> (intersection-over-union threshold of 0.5) and mAP<sub>50-95</sub> (averaged over thresholds from 0.5 to 0.95). The prior work at UBM is documented in the thesis by [Fusa \(2025\)](#), which established the YOLOv11n baseline and the stereo matching pipeline we build upon. That thesis reports mAP<sub>50</sub>=0.824 on the FSOCO validation set, though as we

discuss in chapter 4, we were unable to reproduce this figure with the available training configuration.

# Chapter 3

## Methods

### 3.1 Experimental Setup

Github makes software production on concurrent machines easy, we used a M1 Macbook Air to iterate on the existing inference pipeline: 'ubm-yolo-detector'. We also used an Ubuntu workstation with a Nvidia 4080 Super for development inside a new repository: 'ml4cv-project'. The existing inference pipeline, 'ubm-yolo-detector', enforced the decision to use Ultralytics models within the ultralytics library, as we considered it important to slot the new model in the existing modular framework. Ultralytics facilitated training, validation, and export and should remain part of the framework. For visualization, we used Weights and Biases (wandb) ([Weights & Biases, Inc., 2024](#)) for experiment tracking and comparison. The Autonomous System Unit (ASU) listens to control commands, sends them to the racecar's control unit and listens for mission parameters. It is a water-cooled desktop station running Ubuntu OS with the Robot Operating System 2 (ROS2) framework and has a 500w power budget ([Fusa, 2025](#)). For our purposes, we care only about its GPU, a Nvidia 4060.

Table 3.1: Hardware used for training and deployment.

Machine	Purpose	GPU	OS
M1 MacBook	Code development	CPU only	macOS
Ubuntu Workstation	Training & testing	RTX 4080 Super (CUDA)	Ubuntu
ASU (Race Car)	Deployment	RTX 4060	Ubuntu (ROS2)

### 3.2 Datasets

We used three datasets, summarized in table 3.2. FSOCO-12 ([FMDV, 2024](#)) is the primary training and benchmarking dataset. The cone-detector dataset ([FSB Driverless, 2024](#)) is a larger collection used exclusively for pre-training in the two-stage strategy. Finally,

fsoco-ubm is a small test set we created from the car’s own camera to validate whether benchmark results translate to the real world.

Table 3.2: Datasets used in this project.

Dataset	Images	Instances	Split	Purpose
FSOCO-12 (train)	7,120	~78,000	Train	Primary training
FSOCO-12 (val)	1,968	~36,000	Validation	Model selection
FSOCO-12 (test)	689	12,054	Test	Standard benchmark
cone-detector	22,725	~200,000	Train	Stage 1 pre-training
fsoco-ubm	96	1,426	Test	Real-world validation

### 3.2.1 FSOCO-12

FSOCO-12 (Formula Student Objects in Context, version 12) is a community-curated dataset hosted on Roboflow (FMDV, 2024; FSOCO Contributors, 2020). It contains 9,777 images split into train (7,120), validation (1,968), and test (689) sets, with approximately 126,000 cone instances across the five FSAE classes. The images come from multiple Formula Student teams and cover diverse lighting, weather, and track conditions at variable resolutions. The class distribution is imbalanced: yellow and blue cones dominate (~77% of test instances), while unknown cones account for only 5.6% and are the hardest to detect. We use the test set (689 images, 12,054 instances) as our primary benchmark throughout this report.




Figure 2: FSOCO supports five object classes. The four main classes are shown here, the fifth class *other* includes all cones that are not rules compliant.

Figure 3.1: Class distribution in the FSOCO-12 dataset.

### 3.2.2 cone-detector

The cone-detector dataset (FSB Driverless, 2024) contains 22,725 images with the same five cone classes as FSOCO-12, making it 2.3× larger. It is also a community dataset

from Roboflow, contributed by a different Formula Student team. Training a model from scratch on this dataset alone plateaus at  $mAP_{50} \approx 0.68$ , which is worse than our FSOCO-12 baseline. However, the volume of data makes it useful for pre-training: the model can learn general cone features from a larger and more varied pool before fine-tuning on the curated FSOCO-12 set. We used it exclusively for Stage 1 of the two-stage training strategy described in section 3.6.

### 3.2.3 fsoco-ubm: Real-World Test Set

FSOCO-12 is an internet dataset assembled from many teams and cameras, so good performance on it does not guarantee good performance on our car. To address this, we created fsoco-ubm: a 96-image test set extracted directly from the ZED 2i stereo camera on the UBM race car. The images were recorded on November 20, 2025 at the Rioveggio test track during driving runs at 30–50 km/h. We extracted one frame every 60 frames from the stereo video (one sample every two seconds of real-world time), split the  $2560 \times 720$  stereo pairs into left and right images of  $1280 \times 720$ , and annotated them using Roboflow Label Assist with manual review.

The dataset contains 1,426 cone instances across the five classes. The distribution differs from FSOCO-12: yellow (53.2%) and blue (40.4%) cones dominate, large orange cones appear in small numbers (4.6%), and orange cones are nearly absent (0.6%). The real-world conditions introduce challenges not present in the internet benchmark: motion blur from the moving car, variable outdoor lighting with shadows and bright sky, small pixel area for distant cones, and color desaturation that makes cone classification harder even when the cone is detected. We use fsoco-ubm strictly as a held-out test set and never for training. Future work expanding on this dataset is encouraged.

## 3.3 Baseline Reproduction

Our first step was to reproduce the baseline from Fusa’s thesis (Fusa, 2025), which reports  $mAP_{50}=0.824$  for YOLOv11n on FSOCO. We trained YOLOv11n for 300 epochs on FSOCO-12 using Ultralytics default hyperparameters (AdamW optimizer,  $lr_0=0.01$ , batch 64,  $640 \times 640$  input, default augmentation with  $\text{mosaic}=1.0$ ,  $\text{mixup}=0.0$ ). The model converged around epoch 200 and plateaued at  $mAP_{50}=0.714$  on the validation set, 13.4% below the thesis claim.

We contacted the thesis author, who confirmed that the reported results were produced by other team members (Gabriele and Patta) using an unknown “particular dataset” and training configuration that was subsequently lost. The production model currently deployed on the car, which we evaluated independently on the FSOCO-12 test set, achieves only  $mAP_{50}=0.666$ . Our baseline therefore already exceeds the production model by 6.2%,

and we adopt  $\text{mAP}_{50}=0.707$  (test set) as the reproducible reference point for all subsequent experiments.

### 3.4 Hyperparameter Sweep

To determine whether the baseline could be improved through hyperparameter tuning, we ran a Bayesian optimization sweep using Weights & Biases ([Weights & Biases, Inc., 2024; Snoek et al., 2012](#)). The sweep explored 13 hyperparameters: learning rate ( $\text{lr}_0$ ,  $\text{lrf}$ ), momentum, weight decay, warmup epochs, close mosaic epoch, dropout, and six augmentation parameters ( $\text{hsv\_h}$ ,  $\text{hsv\_s}$ ,  $\text{hsv\_v}$ ,  $\text{mosaic}$ ,  $\text{mixup}$ ,  $\text{copy\_paste}$ ). Each run trained for 100 epochs on FSOCO-12 with YOLOv11n.

Of 21 planned runs, 10 completed successfully and 11 crashed. The crashes were strongly correlated with aggressive augmentation: crashed runs had  $4\times$  higher mixup (0.197 vs 0.049) and nearly  $2\times$  higher dropout (0.156 vs 0.081) on average. High mixup creates heavily blended training images that, combined with high dropout, introduce too much regularization for the model to converge. Among the 10 completed runs, the best achieved  $\text{mAP}_{50}=0.709$ , which is 0.7% *worse* than our baseline. The mean across all runs was 0.703, with a standard deviation of only 0.019 (2.7% of baseline). No run improved over the defaults. This narrow variance demonstrates that YOLOv11n on FSOCO-12 is largely insensitive to these hyperparameters, and that the Ultralytics defaults—which disable both mixup and dropout by default—are already near-optimal. While this sweep does not constitute a formal hypothesis test, the absence of any improvement and the low variance across runs provide strong evidence that further gains are unlikely to come from hyperparameter tuning alone. We stopped the sweep and pivoted to architecture changes.

### 3.5 Architecture Evaluation

Since hyperparameter tuning proved ineffective, we evaluated whether newer architectures could break through the YOLOv11n performance ceiling. We trained YOLO12n ([Tian et al., 2025](#)) and YOLO26n ([Ultralytics, 2025](#)) under identical conditions: FSOCO-12 dataset, 300 epochs, batch 64, Ultralytics default hyperparameters. All three models are nano variants with comparable parameter counts (2.51–2.59M) and GFLOPs (5.8–6.4), as shown in table 3.3, making the comparison fair with respect to model capacity. YOLO12n finished at  $\text{mAP}_{50}=0.708$  on the test set, a marginal gain over YOLOv11n (0.707). YOLO26n reached  $\text{mAP}_{50}=0.763$ , a substantial 7.9% improvement over our baseline and 14.6% over the UBM production model. Every class improved, with the largest absolute gains on unknown cones (+23.4%) and orange cones (+6.8%). The results confirmed that for this task, architecture matters more than hyperparameter tuning: YOLO26n with default parameters outperformed every sweep run by a wide margin.

Table 3.3: Architecture specifications for the three evaluated models.

Model	Params (M)	GFLOPs	Year
YOLOv11n	2.59	6.4	2024
YOLO12n	2.56	6.3	2025
YOLO26n	2.51	5.8	2025

## 3.6 Two-Stage Training Strategy

Given that YOLO26n achieved the best single-stage result, we investigated whether pre-training on a larger dataset could push it further. The cone-detector dataset ([FSB Driverless, 2024](#)) contains 22,725 images with the same five classes as FSOCO-12, providing  $2.3\times$  more training data. Both datasets target the same task—cone detection—so the domain gap is small, which makes transfer learning viable. The hypothesis was that the model would learn more robust cone features from the larger pool and then refine them on the curated FSOCO-12 distribution.

Stage 1 pre-trained YOLO26n on cone-detector starting from COCO-pretrained weights. Training was planned for 400 epochs but the loss converged early and we stopped at epoch 338, achieving  $mAP_{50}=0.734$  on the cone-detector validation set. Stage 2 fine-tuned the Stage 1 checkpoint on FSOCO-12 for 300 additional epochs. Our first attempt at Stage 2 failed: the model’s  $mAP_{50}$  dropped from 0.754 to 0.628 in just two epochs—catastrophic forgetting ([French, 1999](#)). The root cause was Ultralytics’ `optimizer='auto'` default, which silently replaced our specified learning rate of 0.001 with a hardcoded 0.01, as detailed in section [A.3](#).

We redesigned Stage 2 as a two-phase process using explicit AdamW ([Loshchilov and Hutter, 2019](#)). Phase 2A froze the first 10 backbone layers and trained only the detection head for 50 epochs with  $lr=0.001$  and 20% warmup (10 epochs), allowing the head to adapt to the FSOCO-12 distribution without disturbing the pre-trained features. Phase 2B then unfroze all layers for 250 epochs with an ultra-low learning rate of 0.00005 (100 $\times$  lower than training from scratch) and 50 epochs of warmup, using cosine annealing for a smooth decay. This eliminated the forgetting entirely: validation  $mAP_{50}$  improved monotonically throughout Stage 2 and converged at 0.761 on the FSOCO-12 test set, comparable to the single-stage result (0.763) but with better recall and real-world precision, as we discuss in chapter [4](#). The results however, do not encourage two-stage retraining due to the modest performance increase tradeoff with much longer computational effort.

## 3.7 Deployment Pipeline

The trained PyTorch model is deployed through a three-step export pipeline. First, the `best.pt` checkpoint is exported to ONNX ([ONNX Community, 2019](#)) with a fixed batch

size of 2 (two images in a stereo pair). Then, on the ASU itself, the ONNX graph is compiled into a TensorRT FP16 engine (NVIDIA Corporation, 2024) using `trtexec`. Building the engine on the target hardware is mandatory because TensorRT optimizes kernel selection and memory layout for the specific GPU architecture (RTX 4060).

The C++ ROS2 inference node (`ros_yolo_detector_node`) receives stereo image pairs from the ZED 2i camera topic and runs YOLO inference on the left and right frames in parallel using two threads. The resulting bounding box vectors are sorted by center coordinates and matched across the stereo pair using a distance and similarity criterion, followed by feature matching for robust association. Matched detections are triangulated into 3D positions in the camera coordinate frame and then transformed into the vehicle frame (FLU convention) for downstream SLAM consumption. The node publishes both an annotated image with bounding box overlays and a `DetectionInfosArray` message containing 3D coordinates and confidence scores. We lowered the detection confidence threshold from the default 0.5 to 0.35, which improved mAP<sub>50</sub> on fsoco-ubm by 10% by recovering additional true positives at the cost of a marginal increase in false detections. This parameter as well as the NMS threshold are modifiable without recompiling and can be tuned further on the track.




Figure 3.2: 3D visualization of detected cones from the stereo pipeline.

# Chapter 4

## Results

We evaluate all models on two test sets: FSOCO-12 (689 images, 12,054 instances), the standard internet benchmark, and fsoco-ubm (96 images, 1,426 instances), a real-world test set from the car’s own camera. The five models compared are: the UBM production YOLOv11n currently deployed on the car, our retrained YOLOv11n baseline, YOLO12n, and YOLO26n trained both single-stage and two-stage. We report  $mAP_{50}$ ,  $mAP_{50-95}$ , precision, and recall following the COCO evaluation protocol (Lin et al., 2014).

### 4.1 FSOCO-12 Benchmark Results

Table 4.1 presents the FSOCO-12 test set results. YOLO26n dominates: the single-stage variant achieves the highest  $mAP_{50}=0.763$  and precision (0.849), while the two-stage variant leads in recall (0.708) and  $mAP_{50-95}=0.528$ . Both YOLO26n variants outperform the UBM production model by over 14%, confirming that the architecture upgrade accounts for the bulk of the improvement. YOLO12n and our retrained YOLOv11n baseline land nearly identically at  $mAP_{50}\approx 0.708$ , both already 6% above the production model. The gap between YOLOv11n and YOLO12n is marginal, while the jump to YOLO26n is substantial.

Table 4.1: Overall performance on the FSOCO-12 test set (689 images, 12,054 instances).

Model	Training Data	Epochs	$mAP_{50}$	$mAP_{50-95}$	Prec.	Recall
YOLO26n (single)	FSOCO-12	300	<b>0.763</b>	0.524	<b>0.849</b>	0.694
YOLO26n (two-stage)	CD→FSOCO-12	338+300	0.761	<b>0.528</b>	0.832	<b>0.708</b>
YOLO12n	FSOCO-12	300	0.708	0.485	0.840	0.654
YOLOv11n (ours)	FSOCO-12	300	0.707	0.490	0.816	0.662
UBM Production	unknown	300	0.666	0.461	0.803	0.579

CD = cone-detector dataset (22,725 images, pre-training stage).

## 4.2 Per-Class Analysis

The per-class breakdown in tables 4.2 and 4.3 reveals a consistent hierarchy across both architectures. Large orange cones are the easiest to detect ( $\text{mAP}_{50}=0.886$  for YOLO26n), which makes sense given their distinctive size and bright color. Blue and yellow cones, the primary track markers and the two most common classes, perform well at 0.863 and 0.856 respectively. Orange cones sit slightly lower at 0.843 despite sharing a color family with large orange, because they are smaller and harder to distinguish from large orange cones. Unknown cones are the hardest class by far at  $\text{mAP}_{50}=0.364$ —a consequence of the class definition itself: these are cones that annotators could not confidently classify, so the model inherits that ambiguity. YOLO26n improves every class over YOLO12n, with the largest absolute gain on unknown cones (+23.4%), suggesting that its architecture better captures the subtle features needed to handle ambiguous cases.

Table 4.2: Per-class performance of YOLO26n (single-stage) on FSOCO-12 test set.

Class	Images	Instances	Prec.	Recall	$\text{mAP}_{50}$	$\text{mAP}_{50-95}$
Large Orange Cone	154	408	0.873	0.833	<b>0.886</b>	0.688
Blue Cone	506	4,437	0.927	0.783	0.863	0.602
Yellow Cone	562	4,844	0.915	0.774	0.856	0.583
Orange Cone	286	1,686	0.892	0.779	0.843	0.571
Unknown Cone	68	679	0.635	0.297	0.364	0.178

Table 4.3: Per-class performance of YOLO12n on FSOCO-12 test set.

Class	Prec.	Recall	$\text{mAP}_{50}$	$\text{mAP}_{50-95}$
Large Orange Cone	0.912	0.821	<b>0.871</b>	0.693
Blue Cone	0.912	0.738	0.804	0.548
Yellow Cone	0.890	0.727	0.796	0.534
Orange Cone	0.879	0.722	0.775	0.525
Unknown Cone	0.607	0.264	0.295	0.124

## 4.3 Two-Stage vs. Single-Stage Training

Table 4.4 compares the two training strategies head-to-head. On the FSOCO-12 benchmark the difference is within noise:  $\text{mAP}_{50}$  differs by just 0.2%, and neither variant is consistently better across all metrics. The two-stage model wins on recall (+1.4 percentage points) and  $\text{mAP}_{50-95}$  (+0.6%), meaning it finds more cones and localizes them more precisely, at the cost of slightly lower confidence scores (precision drops 1.6 percentage points). On the real-world fsoco-ubm set, the two models tie at  $\text{mAP}_{50}=0.565$ , but two-stage achieves

3.4 percentage points higher precision, indicating fewer false positives under deployment conditions. We deployed the two-stage model because in autonomous racing, missing a cone is more dangerous than a false detection: the planning algorithm can filter spurious detections, but a missed cone leaves the car blind to a track boundary. However, we do not recommend training two stage models due to their quadrupled computational cost; better to train other architectures and improve the datasets.

Table 4.4: Two-stage vs. single-stage YOLO26n training comparison.

Metric	Single-Stage	Two-Stage	Delta
FSOCO-12 mAP <sub>50</sub>	<b>0.763</b>	0.761	-0.2%
FSOCO-12 mAP <sub>50-95</sub>	0.524	<b>0.528</b>	+0.6%
FSOCO-12 Precision	<b>0.849</b>	0.832	-1.6 pp
FSOCO-12 Recall	0.694	<b>0.708</b>	+1.4 pp
fsoco-ubm mAP <sub>50</sub>	0.565	<b>0.565</b>	+0.04%
fsoco-ubm Precision	0.615	<b>0.649</b>	+3.4 pp
Generalization Gap	-25.9%	<b>-25.8%</b>	+0.1 pp

pp = percentage points.

## 4.4 Hyperparameter Sweep

The sweep results in table 4.5 confirm that hyperparameter tuning is not a productive direction for this task. The best sweep run ( $mAP_{50}=0.709$ ) fell below our baseline (0.714), and the mean across all completed runs was 0.703 with a standard deviation of only 0.019. The 2.7% spread shows that YOLOv11n on FSOCO-12 is largely insensitive to the hyperparameters we explored, and the Ultralytics defaults sit near the top of the distribution. Ultralytics uses a MuSGD Optimizer enabling more stable training and faster convergence and we verified the stable training claim. (Ultralytics, 2025)

Table 4.5: Hyperparameter sweep summary (W&B Bayesian optimization).

Metric	Value
Runs completed	10 / 21
Best sweep mAP <sub>50</sub>	0.709
Baseline mAP <sub>50</sub>	0.714
Mean of sweep runs	0.703
Standard deviation	0.019

## 4.5 Real-World Validation (fsoco-ubm)

Table 4.6 shows performance on fsoco-ubm. Every model suffers a substantial drop from the FSOCO-12 benchmark, ranging from  $-21.5\%$  (YOLOv11n) to  $-27.0\%$  (YOLO12n). This gap reflects the harder conditions in real-world data: motion blur from driving at 30–50 km/h, variable outdoor lighting with shadows and bright sky, and small pixel area for distant cones. Despite the drop, the YOLO26n variants maintain first place at  $mAP_{50}=0.565$ . An interesting finding is that YOLOv11n generalizes best, losing only 21.5% and achieving the highest precision on this set (0.874), which suggests that its simpler architecture is more conservative and less prone to false positives on out-of-distribution data. YOLO12n generalizes worst among our trained models ( $-27.0\%$ ) and ends up tied with UBM production at 0.517.

Table 4.6: Real-world performance on fsoco-ubm test set (96 images, 1,426 instances).

Model	$mAP_{50}$	Prec.	Recall	Gap vs FSOCO-12
YOLO26n (two-stage)	<b>0.565</b>	<b>0.649</b>	0.462	$-25.8\%$
YOLO26n (single)	<b>0.565</b>	0.615	<b>0.469</b>	$-25.9\%$
YOLOv11n (ours)	0.555	0.874	0.447	$-21.5\%$
YOLO12n	0.517	0.572	0.454	$-27.0\%$
UBM Production	0.517	0.635	0.393	$-22.3\%$

## 4.6 Deployment Performance

Table 4.7 breaks down the inference latency on the RTX 4060 onboard the race car. YOLO26n completes a forward pass in 2.63 ms on average, of which only 1.02 ms is GPU compute—the remainder is host-to-device memory transfer (1.58 ms) and device-to-host transfer (0.03 ms). The pipeline is transfer-bound, not compute-bound, which means a more powerful GPU would not significantly reduce latency; the bottleneck is moving the image data across the PCIe bus. At 2.63 ms per image, the model has a  $6.3\times$  margin over the 16.7 ms budget required for 60 fps operation, so we did not pursue INT8 quantization: the typical 1–2%  $mAP_{50}$  penalty is not worth the marginal speed gain when the model is already well within the real-time envelope. Table 4.8 summarizes the full upgrade over the previous production system. The deployed YOLO26n is 14.6% more accurate on FSOCO-12, 9.4% more accurate on fsoco-ubm, and runs at 2.63 ms versus the 6.78 ms reported for the previous production model, all while using fewer parameters (2.51M vs 2.59M). The satisfactory performance and 1 millisecond compute are desirable but the latency margin future work leaves room for training a YOLO27s model and even larger variants to quantify the performance vs latency tradeoff.

Table 4.7: Inference latency breakdown on RTX 4060 (TensorRT FP16).

Component	YOLO26n (ms)	YOLOv11n prod. (ms)
H2D Transfer	1.58	1.58
GPU Compute	<b>1.02</b>	0.99
D2H Transfer	0.03	0.13
<b>Total</b>	<b>2.63</b>	2.70
Max FPS	380	370

Table 4.8: Deployment comparison: YOLO26n vs. UBM baseline.

Metric	YOLO26n	UBM Baseline
Architecture	YOLO26n	YOLOv11n
Parameters	2.51M	2.59M
GFLOPs	5.8	6.4
TensorRT FP16 Size	9.35 MB	~9 MB
Mean Latency (RTX 4060)	2.63 ms	6.78 ms
Throughput	633 qps	~147 qps
mAP <sub>50</sub> (FSOCO-12)	0.763	0.666
mAP <sub>50</sub> (fsoco-ubm)	0.565	0.517
Real-time margin (60 fps)	6.3×	2.5×

# Chapter 5

## Conclusion

### 5.1 Summary of Contributions

We systematically evaluated three generations of YOLO architectures for cone detection in Formula Student Driverless. Training YOLOv11n, YOLO12n, and YOLO26n under identical conditions on the FSOCO-12 dataset, we found that YOLO26n—the newest and lightest of the three—achieves the best results:  $mAP_{50}=0.763$  on the standard benchmark, a 14.6% improvement over the model previously deployed on the car. A two-stage training strategy, pre-training on the 22,725-image cone-detector dataset before fine-tuning on FSOCO-12, yielded comparable benchmark accuracy (0.761 vs 0.763) but better recall (+1.4 percentage points) and higher real-world precision (+3.4 percentage points). A Bayesian hyperparameter sweep over 13 parameters demonstrated that the Ultralytics defaults are already near-optimal, and that further gains from tuning alone are unlikely.

To validate our results beyond internet benchmarks, we created fsoco-ubm, a 96-image test set from the car’s own ZED 2i stereo camera at the Rioveggio test track. All models lost 22–27% accuracy on this set compared to FSOCO-12, confirming that curated benchmarks overstate deployment performance. Despite this drop, YOLO26n maintained first place. We deployed the final model via TensorRT FP16 on the onboard RTX 4060, achieving 2.63 ms inference latency—a  $6.3\times$  margin over the 60 fps requirement.

### 5.2 Lessons Learned

Contributing to the existing codebase was straightforward thanks to the modular software architecture established by [Fusa \(2025\)](#): detection, stereo matching, and triangulation are separate stages, so swapping the YOLO model required only replacing the TensorRT engine and updating the confidence threshold. This report also fills a documentation gap. The previous training configuration, hyperparameters, and dataset version were lost when team members graduated, making it impossible to reproduce the originally reported

results. We have documented every training run, dataset version, and export step so that future team members can reproduce and extend this work. The key technical lessons are:

- **Architecture selection matters more than hyperparameter tuning.** The sweep explored 13 parameters across 21 runs and found no improvement over defaults, while upgrading from YOLOv11n to YOLO26n gained 7.9% mAP<sub>50</sub>.
- **Always validate on real data.** FSOCO-12 performance does not reliably predict deployment performance: all models lost 22–27% on our car’s camera data. Without fsoco-ubm, we would have overestimated the deployed model’s accuracy.
- **Two-stage training provides marginal gains at this data scale.** Pre-training on 3× more data improved recall and real-world precision, but the benchmark mAP<sub>50</sub> remained within 0.2% of single-stage. The 43 hours of additional compute are hard to justify for the modest improvement.
- **Verify framework internals.** The Ultralytics `optimizer='auto'` setting silently overwrites user-specified learning rates, which caused catastrophic forgetting during fine-tuning and cost a full day of debugging (section A.3).

## 5.3 Limitations

Several limitations should be noted. We could not reproduce the mAP<sub>50</sub>=0.824 reported in Fusa (2025): the training configuration and dataset version used to produce that figure were lost, and the best we achieved with YOLOv11n on FSOCO-12 under default hyperparameters was 0.707, which we adopted as our reproducible baseline. The fsoco-ubm test set is small (96 images) and captures a single track on a single day, so it does not cover the full range of conditions the car will face at competition—different track layouts, rain, and dusk lighting are absent. The unknown cone class remains poorly detected at mAP<sub>50</sub>=0.364, which is a dataset-level problem: the class is defined by annotator uncertainty rather than a consistent visual feature, so the model inherits that ambiguity. We only evaluated nano-sized models ( $\sim$ 2.5M parameters); larger variants may achieve higher accuracy within the available latency budget but were not tested. Finally, we did not run ablation studies on individual augmentation parameters, so we cannot isolate which strategies (mosaic, mixup, copy-paste) are most effective for cone detection.

## 5.4 Future Work

We recommend the following directions for continued development:

- **Evaluate larger YOLO26 variants.** The  $6.3\times$  latency margin leaves room for the small or medium variants, which may improve detection of distant and ambiguous cones. Even a model  $3\times$  slower would remain real-time.
- **Expand fsoco-ubm.** Add images from different tracks, weather conditions, and times of day. The current 96 images from a single session are a starting point, not a representative benchmark.
- **Target the unknown cone class.** Data augmentation strategies such as color jittering and occlusion simulation may help the model learn to flag uncertain detections. Alternatively, merging unknown cones into the nearest color class during training could improve overall recall at the cost of classification granularity.
- **Auto-exposure compensation.** The ZED 2i camera’s auto-exposure underexposes cones when the sky dominates the frame. A preprocessing step that adjusts exposure based on the lower half of the image could reduce this failure mode.
- **End-to-end integration testing.** We validated detection accuracy and inference speed independently but did not test the full autonomous stack (detection → matching → triangulation → SLAM → planning → control) with the new model under race conditions.

UBM will compete at Formula Student Germany from 11–16 August 2026. The goal is to complete all four dynamic events—Trackdrive, Autocross, Skid Pad, and Acceleration—for the first time; last year a broken axle ended the campaign before we could attempt all of them. The priority is reliability over outright speed: finishing every event and collecting data is more valuable than optimizing for a single fast lap, because the data enables iteration. The improvements in this report contribute to that goal by providing a more accurate and faster perception pipeline, but the real test will come at competition.

# Appendix A

## Technical Details

### A.1 Hyperparameter Sweep Configuration

```
1 # TODO: Paste actual sweep YAML from wandb_sweep.yaml
2 method: bayes
3 metric:
4   name: metrics/mAP50(B)
5   goal: maximize
6 parameters:
7   lr0:
8     min: 0.001
9     max: 0.05
10  momentum:
11    min: 0.8
12    max: 0.98
13  # ... (add remaining 11 parameters)
```

Listing A.1: W&B Bayesian sweep configuration.

### A.2 Two-Stage Training Hyperparameters

### A.3 The `optimizer='auto'` Bug

### A.4 fsoco-ubm Dataset Creation Pipeline

```
1 # TODO: Paste key extract from extract_frames_from_avi.py
2 import cv2
3
4 cap = cv2.VideoCapture("media/lidar1.avi")
```

Table A.1: Two-stage training configuration.

Parameter	Stage 1 (cone-detector)	Stage 2 (FSOCO-12)
Dataset	cone-detector (22,725 imgs)	FSOCO-12 (5,536 imgs)
Epochs	338 (converged early)	300
Batch size	64	64
Optimizer	auto (SGD)	AdamW
Learning rate	0.01	0.001
Image size	640	640
Freeze layers	None	Phase-dependent
Pretrained weights	COCO (yolo26n.pt)	Stage 1 best.pt

```

5 frame_count = 0
6 saved = 0
7 while cap.isOpened():
8     ret, frame = cap.read()
9     if not ret:
10         break
11     if frame_count % 60 == 0: # every 2 seconds
12         cv2.imwrite(f"ubm_test_set/images/lidar1_{saved:04d}.png"
13                     , frame)
14         saved += 1
15 frame_count += 1
cap.release()

```

Listing A.2: Frame extraction from AVI video.

# Bibliography

- Alexey Bochkovskiy, Chien-Yao Wang, and Hong-Yuan Mark Liao. YOLOv4: Optimal speed and accuracy of object detection. *arXiv preprint arXiv:2004.10934*, 2020.
- Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. The PASCAL visual object classes (VOC) challenge. *International Journal of Computer Vision*, 88(2):303–338, 2010. doi: 10.1007/s11263-009-0275-4.
- FMDV. FSOCO cone detection dataset (version 12), 2024. URL <https://universe.roboflow.com/fmdv/fsoco-kxq3s/dataset/12>. Roboflow Universe. Accessed January 2026.
- Formula Student Germany. Formula Student Germany — driverless, 2024. URL <https://www.formulastudent.de/fsg/>. Accessed January 2026.
- Robert M. French. Catastrophic forgetting in connectionist networks. *Trends in Cognitive Sciences*, 3(4):128–135, 1999. doi: 10.1016/S1364-6613(99)01294-2.
- FSB Driverless. Cone detector dataset (version 1), 2024. URL <https://universe.roboflow.com/fsbdriverless/cone-detector-zruok/dataset/1>. 22,725 images. Roboflow Universe. Accessed January 2026.
- FSOCO Contributors. FSOCO: Formula student objects in context, 2020. URL <https://www.fsoco-dataset.com/>. Community dataset for Formula Student cone detection.
- Edoardo Fusa. Pushing cars’ limits: Exploring autonomous technologies in the formula SAE driverless competition. Master’s thesis, University of Bologna, 2025. Stereocamera pipeline for autonomous racing.
- Glenn Jocher, Ayush Chaurasia, and Jing Qiu. Ultralytics YOLOv8, 2023. URL <https://github.com/ultralytics/ultralytics>. Accessed January 2026.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO: Common objects in context. *arXiv preprint arXiv:1405.0312*, 2014.

- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- NVIDIA Corporation. NVIDIA TensorRT: Programmable inference accelerator, 2024. URL <https://developer.nvidia.com/tensorrt>. Accessed January 2026.
- ONNX Community. ONNX: Open neural network exchange, 2019. URL <https://onnx.ai/>. Accessed January 2026.
- Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016.
- Roboflow, Inc. Roboflow: Computer vision tools, 2024. URL <https://roboflow.com/>. Dataset management and annotation. Accessed January 2026.
- SAE International. FSAE driverless supplement v11, 2025. URL <https://www.saeonline.com/CompResources/2025/8f030a58-d9e4-49b8-bc83-6ca16c7ce715/FSAE-Driverless-Supplement-V11.pdf>. Section: Dynamic Event Markings.
- Jasper Snoek, Hugo Larochelle, and Ryan P. Adams. Practical Bayesian optimization of machine learning algorithms. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 25, 2012.
- Stereolabs. ZED 2i stereo camera, 2021. URL <https://www.stereolabs.com/products/zed-2i>. Accessed January 2026.
- Yunjie Tian, Qiang Ye, and David Doermann. YOLO12: Attention-centric real-time object detectors. *arXiv preprint arXiv:2502.12524*, 2025.
- Ultralytics. YOLO11: Real-time object detection, 2024. URL <https://docs.ultralytics.com/models/yolo11/>. Accessed January 2026.
- Ultralytics. YOLOv26: Advancing real-time object detection, 2025. URL <https://docs.ultralytics.com/models/yolo26/>. Accessed January 2026.
- Weights & Biases, Inc. Weights & Biases, 2024. URL <https://wandb.ai/>. Experiment tracking platform. Accessed January 2026.
- Hongyi Zhang, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. *arXiv preprint arXiv:1710.09412*, 2018.